

COMP90042 LECTURE 4

SEQUENCE TAGGING: HIDDEN MARKOV MODELS

OVERVIEW

- ▶ Tagging sequences
 - ▶ modelling concepts
 - ▶ Markov chains and hidden Markov models (HMMs)
 - ▶ decoding: finding best tag sequence
- ▶ Application to POS tagging

SEQUENTIAL PREDICTION

- ▶ Aim to predict *sequence* of labels for a sentence
 - ▶ instance of more general '*structured prediction*' which includes parsing
- ▶ Tagging common in NLP, e.g., part of speech
 - ▶ **Input:** *I see a silhouette of a man.*
 - ▶ **Output:** *PRP VBZ DT NN PP DT NN .*
- ▶ Other popular sequence tagging tasks include *named entity*, *shallow parsing (chunking)*
- ▶ How to build a classifier over sequences, which might be of differing lengths?

NAÏVE APPROACHES FOR SEQ. PRED.

1. Treat as one big classification label:

- e.g., $\mathbf{t} = PRP_VBZ_DT_NN_PP_DT_NN_$.
- but there are exponentially many combinations,
 $|\text{Tags}|^M$ for input of length M (too many parameters!)
- and how to tag sequences of differing lengths?

2. As *independent* classification problems

- ***I** see a silhouette of a man.* $\rightarrow t_1 = PRP$
- *I **see** a silhouette of a man.* $\rightarrow t_2 = VBZ \quad \dots$
- much simpler model & # parameters
- most useful information are the neighbouring tags
(but *can't be used*)

Notation key:


Bold (\mathbf{t}) means *vector*
of several values

Normal (t) means
one value

Capital (A) means
table (matrix) of
values

A BETTER APPROACH TO SEQ. PRED.

- A solution: learning a local classifier
 - e.g., $\Pr(t_n \mid w_n, t_{n-2}, t_{n-1})$ or $P(w_n, t_n \mid t_{n-1})$
 - a more practical model, has key local information available




- But raises *search* problem: how to find the best tag sequence for each word
 - can we avoid the exponential complexity in search?

FEATURES FOR POS TAGGING

- ▶ For part-of-speech tagging, most important features are:
 - ▶ current word
 - ▶ tags for adjacent words (*why just to the left?*)
- ▶ E.g., for *Time **flies** like an arrow*;
 - ▶ *flies* mostly seen as either a VERB or NOUN
 - ▶ *flies* is likely a VERB if previous tag is NOUN
 - ▶ *flies* is likely a NOUN if previous tag is an ADJECTIVE



Time Flies Like
An Arrow



Fruit Flies Like
A Banana

<http://likesuccess.com/img4568782>

MARKOV CHAINS

- ▶ Useful trick to decompose complex chain of events into simpler, smaller modellable events; e.g.,
 - ▶
$$\Pr(t_1 t_2 \dots t_n \mid w_1 w_2 \dots w_n) = \Pr(t_1 \mid w_1) \Pr(t_2 \mid w_2 t_1) \dots \Pr(t_n \mid w_n t_{n-1})$$
 MEMM
 - ▶
$$\Pr(t_1 t_2 \dots t_n w_1 w_2 \dots w_n) = \Pr(t_1 w_1) \Pr(t_2 w_2 \mid t_1) \dots \Pr(t_n w_n \mid t_{n-1})$$
 HMM
- ▶ Make some simplifying assumptions
 - ▶ *Markov assumption*: Only fixed number k of recent tags are relevant (k is known as the *Markov order*; in above $k=1$)
 - ▶ *Limited dependency between words and their tags*: Tags are assumed to capture the local context needed to explain the observations

MARKOV MODELS

- Characterised by
 - set of states
 - initial state occ prob
 - state transition probs
 - outgoing edges normalised
- Can score sequences of observations
 - For stock price example:
up-up-down-up-up
 - maps directly to states
 - simply multiply probs

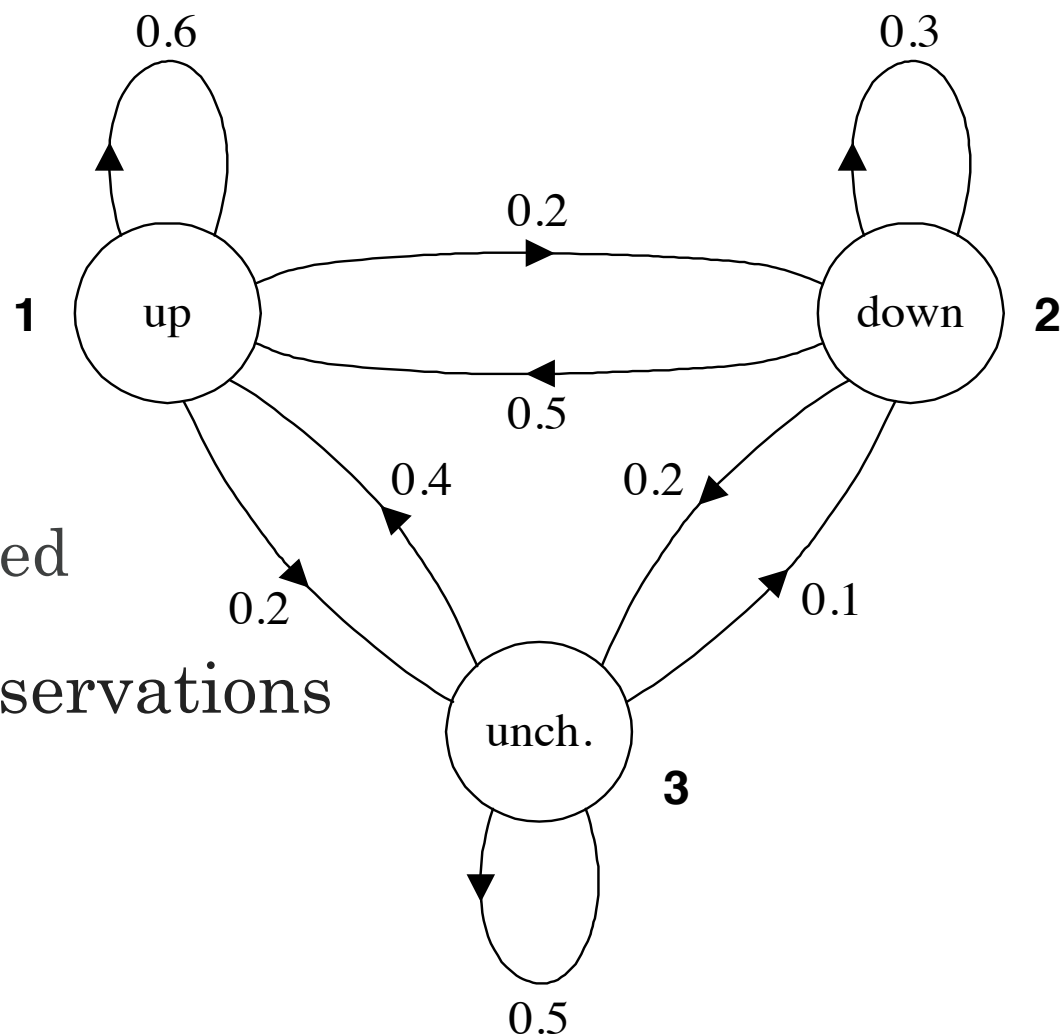


Fig. from Spoken language processing; Huang, Acero, Hon (2001); Prentice Hall

HIDDEN MARKOV MODELS

- ▶ Each state now has in addition

- ▶ emission prob vector

- ▶ No longer 1:1 mapping $\begin{bmatrix} 0.7 \\ 0.1 \\ 0.2 \end{bmatrix}$

- ▶ from observation sequence to states

- ▶ E.g., up-up-down-up-up could be generated from any state sequence

- ▶ but some more likely than others!

- ▶ State sequence is ‘hidden’

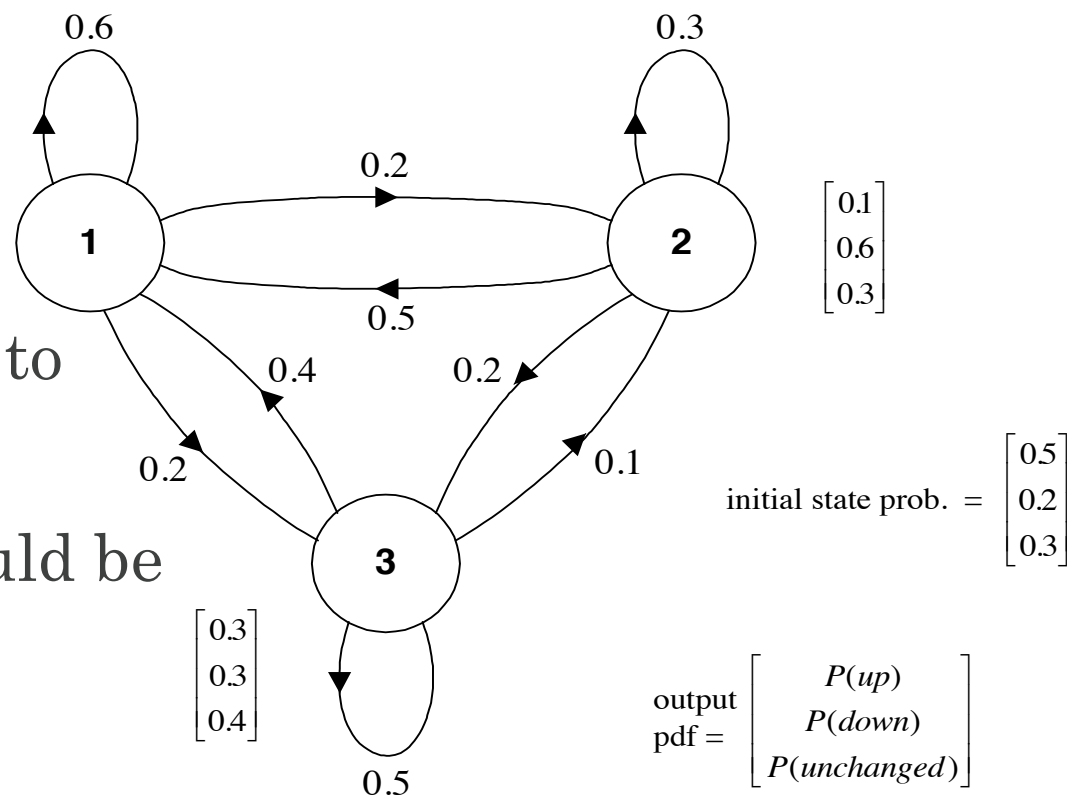


Fig. from Spoken language processing; Huang, Acero, Hon (2001); Prentice Hall

NOTATION

- Basic units are a sequence of
 - O , observations e.g., words
 - Ω , states e.g., POS tags
- Model characterised by
 - initial state probs π = vector of $|\Omega|$ elements
 - transition probs A = matrix of $|\Omega| \times |\Omega|$
 - emission probs O = matrix of $|\Omega| \times |O|$
- Together define the probability of a sequence
 - of observations together with their tags
 - a model of $P(w, t)$
- Notation: w = observations; t = tags; i = time index

ASSUMPTIONS

- Two assumptions underlying the HMM
- Markov assumption
 - states independent of all but most recent state
 - $P(t_i \mid t_1, t_2, t_3, \dots, t_{i-2}, t_{i-1}) = P(t_i \mid t_{i-1})$
 - i.e., state sequence is a *Markov chain*
- Output independence
 - outputs dependent only on matching state
 - $P(w_i \mid w_1, t_1, \dots, w_{i-1}, t_{i-1}, t_i) = P(w_i \mid t_i)$
 - forces the state t_i to carry all information linking w_i with neighbours
- Are these assumptions realistic?

PROBABILITY OF SEQUENCE

- Probability of sequence “up-up-down”

State seq	up π	O	up A	O	down A	O	total
1,1,1	0.5 x	0.7 x	0.6 x	0.7 x	0.6 x	0.1 =	0.00882
1,1,2	0.5 x	0.7 x	0.6 x	0.7 x	0.2 x	0.6 =	0.01764
1,1,3	0.5 x	0.7 x	0.6 x	0.7 x	0.2 x	0.3 =	0.00882
1,2,1	0.5 x	0.7 x	0.2 x	0.1 x	0.5 x	0.1 =	0.00035
1,2,2	0.5 x	0.7 x	0.2 x	0.1 x	0.3 x	0.6 =	0.00126
...							
3,3,3	0.3 x	0.3 x	0.5 x	0.3 x	0.5 x	0.3 =	0.00203

– 1,1,2 is the highest prob hidden sequence

– total prob is 0.054398, not 1 ... why??

HMM DECODING PROBLEM

- ▶ Given observation sequence(s)
 - ▶ e.g., up-up-down-up-up
- ▶ **Raises inference problem**
 - ▶ what states were used to create this sequence?
 - ▶ *Viterbi* algorithm solves for the most likely states, also called '*decoding*'
- ▶ Other key inference problem (**not covered here!**)
 - ▶ unsupervised estimation where training labels are *hidden*, uses variant of the Expectation Maximisation (EM) algorithm

HMMS FOR TAGGING

- Recall part-of-speech tagging
 - time/**Noun** flies/**Verb** like/**Prep** an/**Art** arrow/**Noun**
- What are the units?
 - words = observations
 - tags = states
- Key challenges
 - estimate model from state-supervised data
e.g., based on frequencies
 - decoding for full tag sequences



EXAMPLE

- time/**Noun** flies/**Verb** like/**Prep** an/**Art** arrow/**Noun**
 - Prob = $P(\text{Noun}) P(\text{time} \mid \text{Noun}) \times$
 $P(\text{Verb} \mid \text{Noun}) P(\text{flies} \mid \text{Verb}) \times$
 $P(\text{Prep} \mid \text{Verb}) P(\text{like} \mid \text{Prep}) \times$
 $P(\text{Art} \mid \text{Prep}) P(\text{an} \mid \text{Art}) \times$
 $P(\text{Noun} \mid \text{Art}) P(\text{arrow} \mid \text{Noun})$
- time/**Noun** flies/**Noun** like/**Verb** an/**Art** arrow/**Noun**
 - Prob = $P(\text{Noun}) P(\text{time} \mid \text{Noun}) \times$
 $P(\text{Noun} \mid \text{Noun}) P(\text{flies} \mid \text{Noun}) \times$
 $P(\text{Verb} \mid \text{Noun}) P(\text{like} \mid \text{Verb}) \times$
 $P(\text{Art} \mid \text{Prep}) P(\text{an} \mid \text{Art}) \times$
 $P(\text{Noun} \mid \text{Art}) P(\text{arrow} \mid \text{Noun})$
- Which do you think is more likely?
- What does a state of the art tagger choose?
<http://nlp.stanford.edu:8080/corenlp/process>

ESTIMATING A VISIBLE MARKOV TAGGER

- Estimation

- what values to use for $P(w \mid t)$?
- what values to use for $P(t_i \mid t_{i-1})$ and $P(t_1)$?
- use simple frequencies over training corpus, i.e.,

$$P(t_i | t_{i-1}) = \frac{c(t_{i-1}, t_i)}{c(t_{i-1})} \qquad P(w_i | t_i) = \frac{c(w_i, t_i)}{c(t_i)}$$

E.g.,

$A_{\text{verb}, \text{noun}}$ = how often does Verb follow Noun,
versus other tags?

$O_{\text{noun}, \text{'flies'}}$ = how often are Nouns written as “flies”,
versus other word types?

– (probably want to smooth counts, e.g., adding 0.5)

PREDICTION

- Prediction
 - given a sentence, \mathbf{w} , find the sequence of tags, \mathbf{t}

$$\begin{aligned}\arg \max_{\mathbf{t}} P(\mathbf{w}, \mathbf{t}) &= P(t_1)P(w_1|t_1) \prod_{i=2}^M P(t_i|t_{i-1})P(w_i|t_i) \\ &= \pi_{t_1} O_{t_1, w_1} \prod_{i=2}^M A_{t_{i-1}, t_i} O_{t_i, w_i}\end{aligned}$$

- problems
 - exponential number of values of \mathbf{t}
 - but computation can be factorised...

VITERBI ALGORITHM

- Form of **dynamic programming** to solve maximisation
 - define matrix α of size M (length) x T (tags)

$$\alpha[i, t_i] = \max_{t_1 \cdots t_{i-1}} P(w_1 \cdots w_i, t_1 \cdots t_i)$$

- full sequence max is then

$$\max_{\vec{t}} P(\vec{w}, \vec{t}) = \max_{t_M} \alpha[M, t_M]$$

- how to compute α ?
- We're interested in the **tags**, not just the **max** value
 - can be recovered from α using back-pointers

VITERBI RECURSION

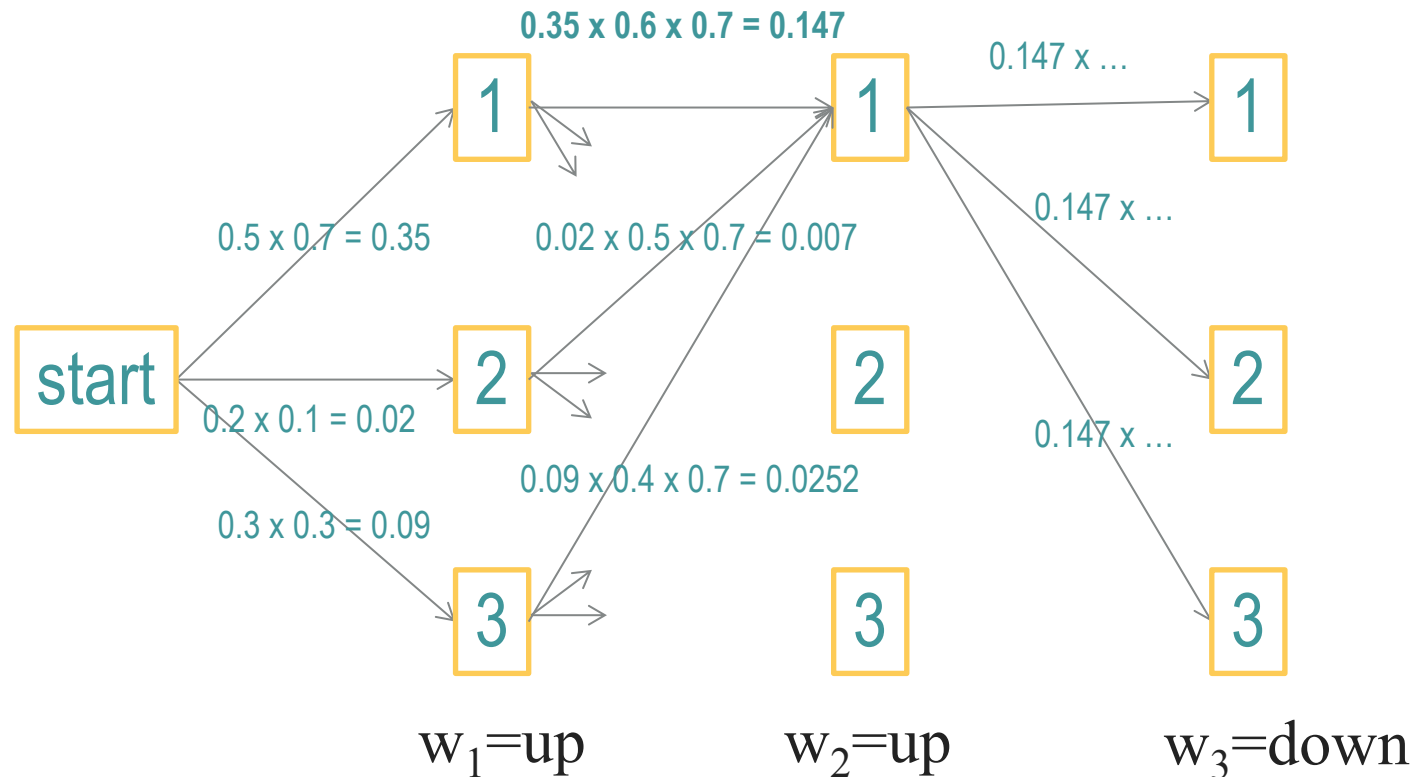
- ▶ Can be defined recursively

$$\begin{aligned}\alpha[i, t_i] &= \max_{t_1 \cdots t_{i-1}} P(w_1 \cdots w_i, t_1 \cdots t_i) \\ &= \max_{t_1} \cdots \max_{t_{i-2}} \max_{t_{i-1}} P(w_1 \cdots w_i, t_1 \cdots t_i) \\ &= \max_{t_1} \cdots \max_{t_{i-2}} \max_{t_{i-1}} P(w_1 \cdots w_{i-1}, t_1 \cdots t_{i-1}) P(w_i, t_i | t_{i-1}) \\ &= \max_{t_{i-1}} \alpha[i-1, t_{i-1}] P(w_i, t_i | t_{i-1})\end{aligned}$$

- ▶ Need a base case to terminate recursion

$$\alpha[1, t_1] = P(w_1, t_1)$$

VITERBI ILLUSTRATION



- ▶ All maximising sequences with $t_2=1$ must also have $t_1=1$
- ▶ No need to consider extending $[2,1]$ or $[3,1]$.

VITERBI ANALYSIS

- ▶ Algorithm as follows

```
alpha = np.zeros(M, T)
for t in range(T):
    alpha[1, t] = pi[t] * O[w[1], t]

for i in range(2, M):
    for t_i in range(T):
        for t_last in range(T):      # t_last means t_{i-1}
            alpha[i, t_i] = np.max(
                alpha[i, t_i],
                alpha[i-1, t_last] * A[t_last, t_i] * O[w[i], t_i])

best = np.max(alpha[M-1, :])
```

- ▶ Time complexity is $O(M T^2)$

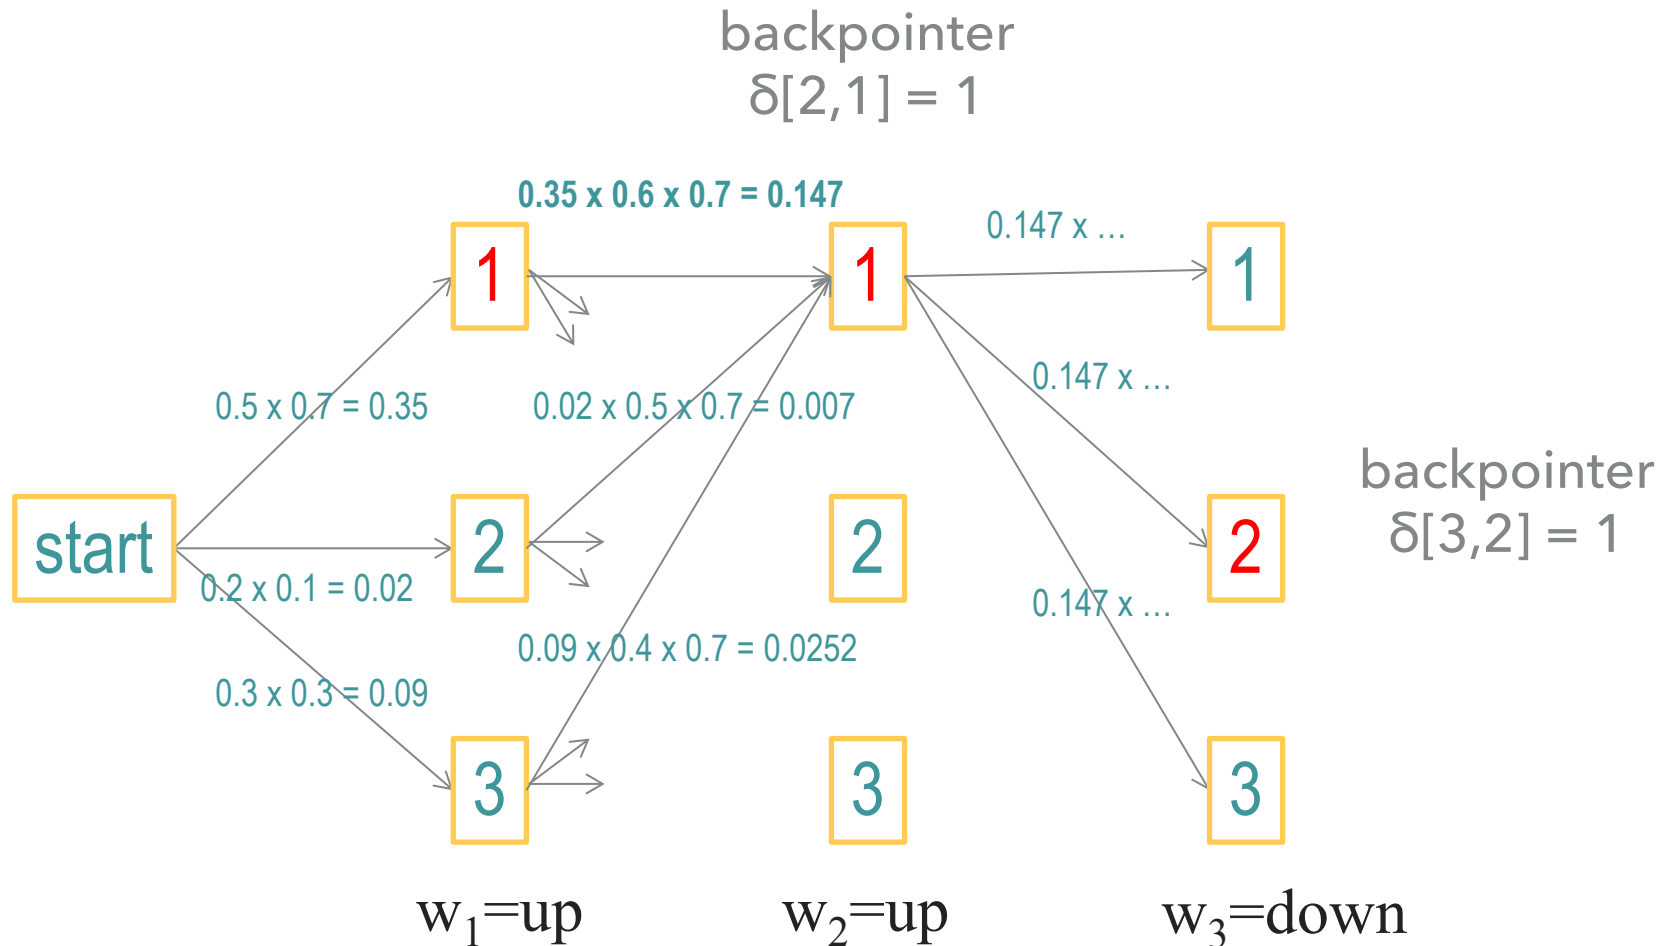
BACKPOINTERS

- Know the maximum score, but not the best path of states
- Solution: don't just store *max* values, α , but who '*won*' each maximisation (the *arg max*), i.e.,

```
for i in range(2, M):
    for t_i in range(T):
        for t_last in range(T):      # t_last means t_{i-1}
            s = alpha[i-1, t_last] * A[t_last, t_i] * O[w[i], t_i]
            if s >= alpha[i, t_i]:
                alpha[i, t_i] = s
                back[i, t_i] = t_last      # record best so far
```

- At the end, must travel backwards from last tag to first to read off the reverse tag sequence

BACKPOINTER ILLUSTRATION



OTHER VARIANT TAGGERS

- HMM is **generative**, $P(t, w)$, ‘creates’ the input
 - allows for unsupervised HMMs: learn model without any tagged data!
- **Discriminative** models also popular, modelling $P(t \mid w)$ directly
 - supports richer feature set, generally better accuracy when trained over large supervised datasets
 - E.g., Maximum Entropy Markov Model (MEMM), Conditional random field (CRF), Connectionist Temporal Classification (CTC)
 - Most *deep learning* models of sequences are discriminative (e.g., encoder-decoders for translation), similar to an MEMM

HMMS IN NLP

- HMMS are highly effective for part-of-speech tagging
 - trigram HMM gets 96.5% accuracy (TnT)
 - related models are state of the art
 - MEMMs 97.3%
 - CRFs 97.6%
 - *English Penn Treebank* tagging accuracy
[https://aclweb.org/aclwiki/index.php?title=POS_Tagging_\(State_of_the_art\)](https://aclweb.org/aclwiki/index.php?title=POS_Tagging_(State_of_the_art))
- Other sequence labelling tasks
 - named entity recognition, shallow parsing, alignment ...
 - In other fields: DNA, protein sequences, image lattices...

SUMMARY

- ▶ Probabilistic models of sequences
- ▶ Introduced hidden Markov models
 - ▶ supervised estimation for learning
 - ▶ Viterbi algorithm for efficient prediction
 - ▶ relationship to other discriminative sequence models

READINGS

- ▶ Hidden Markov models
 - ▶ Jurafsky & Martin 2nd Ed., chapter 6
- ▶ [Optional] Rabiner's HMM tutorial, for more details
 - ▶ <http://tinyurl.com/2hqaf8>
- ▶ **[Just for fun!]** Contemporary sequence tagging methods
 - ▶ Lafferty et al, Conditional random fields: Probabilistic models for segmenting and labeling sequence data (2001), ICML