

COMP90042 LECTURE 6

SEQUENCE TAGGING: HIDDEN MARKOV MODELS

POS TAGGING RECAP

- ▶ Janet will back the bill
- ▶ Janet/**NNP** will/**MB** back/**VP** the/**DT** bill/**NN**
- ▶ Local classifier: prone to **error propagation**
- ▶ What about treating the full sequence as a “class”?
 - ▶ Output: “NNP_MB_VP_DT_NN”
- ▶ Problems:
 - ▶ Exponentially many combinations: $|\text{Tags}|^M$, where M is the length
 - ▶ How to tag sequences of different lengths?

A BETTER APPROACH

- ▶ Tagging is a sentence-level task but as humans we **decompose** it into small word-level tasks.
 - ▶ Janet/**NNP** will/**MD** back/**VP** the/**DT** bill/**NN**
- ▶ Solution:
 - ▶ Define a model that decompose a tagging sentence into individual words
 - ▶ But that takes into account the whole sequence when learning and predicting (no error propagation)
- ▶ This is the idea of **sequence labelling**, and more general, **structured prediction**.

A PROBABILISTIC MODEL

- ▶ Goal: obtain the best tag sequence \mathbf{t} from a sentence \mathbf{w}
 - ▶ $\hat{\mathbf{t}} = \operatorname{argmax}_{\mathbf{t}} P(\mathbf{t}|\mathbf{w})$
 - ▶ $\hat{\mathbf{t}} = \operatorname{argmax}_{\mathbf{t}} \frac{P(\mathbf{w}|\mathbf{t})P(\mathbf{t})}{P(\mathbf{w})} = \operatorname{argmax}_{\mathbf{t}} P(\mathbf{w}|\mathbf{t}) P(\mathbf{t})$ [Bayes]
- ▶ Let's decompose:
 - ▶ $P(\mathbf{w}|\mathbf{t}) = \prod_{i=1}^n P(w_i|t_i)$ [Prob. of a word depends only on the tag]
 - ▶ $P(\mathbf{t}) = \prod_{i=1}^n P(t_i|t_{i-1})$ [Prob. of a tag depends only on the previous tag]
- ▶ These are **independence assumptions** (remember Naïve Bayes?)
- ▶ This is a Hidden Markov Model (HMM)

HIDDEN MARKOV MODEL

- ▶ $\hat{t} = \operatorname{argmax}_t P(\mathbf{w}|\mathbf{t}) P(\mathbf{t})$
- ▶ $P(\mathbf{w}|\mathbf{t}) = \prod_{i=1}^n P(w_i|t_i)$
- ▶ $P(\mathbf{t}) = \prod_{i=1}^n P(t_i|t_{i-1})$
- ▶ Why “Markov”?
 - ▶ Because it assumes the sequence follows a Markov chain: probability of an event (tag) depends only on the previous one (previous tag)
- ▶ Why “Hidden”?
 - ▶ Because the events (tags) are not seen: goal is to find the best sequence

HMMS - TRAINING

- ▶ Parameters are the individual probabilities $P(w_i|t_i)$ and $P(t_i|t_{i-1})$
 - ▶ Respectively, **emission** and **transition** probabilities
- ▶ Training uses Maximum Likelihood Estimation (MLE)
 - ▶ In Naïve Bayes, this is done by simply counting word frequencies according to the class.
- ▶ We do **exactly the same** in HMMs!
 - ▶ $P(\textit{like}|\textit{VB}) = \frac{\textit{count}(\textit{VB},\textit{like})}{\textit{count}(\textit{VB})}$
 - ▶ $P(\textit{NN}|\textit{DT}) = \frac{\textit{count}(\textit{DT},\textit{NN})}{\textit{count}(\textit{DT})}$

HMMS - TRAINING

- ▶ What about the first tag?
- ▶ Assume we have a symbol “<s>” that represents the start of your sentence.
 - ▶
$$P(NN | <s>) = \frac{\text{count}(<s>, NN)}{\text{count}(<s>)}$$
- ▶ What about unseen (word,tag) and (tag, previous) combinations?
- ▶ Same as Naïve Bayes
 - ▶ Smoothing techniques

TRANSITION MATRIX

	NNP	MD	VB	JJ	NN	RB	DT
<s>	0.2767	0.0006	0.0031	0.0453	0.0449	0.0510	0.2026
NNP	0.3777	0.0110	0.0009	0.0084	0.0584	0.0090	0.0025
MD	0.0008	0.0002	0.7968	0.0005	0.0008	0.1698	0.0041
VB	0.0322	0.0005	0.0050	0.0837	0.0615	0.0514	0.2231
JJ	0.0366	0.0004	0.0001	0.0733	0.4509	0.0036	0.0036
NN	0.0096	0.0176	0.0014	0.0086	0.1216	0.0177	0.0068
RB	0.0068	0.0102	0.1011	0.1012	0.0120	0.0728	0.0479
DT	0.1147	0.0021	0.0002	0.2157	0.4744	0.0102	0.0017

Figure 10.5 The A transition probabilities $P(t_i|t_{i-1})$ computed from the WSJ corpus without smoothing. Rows are labeled with the conditioning event; thus $P(VB|MD)$ is 0.7968.

EMISSION (OBSERVATION) MATRIX

	Janet	will	back	the	bill
NNP	0.000032	0	0	0.000048	0
MD	0	0.308431	0	0	0
VB	0	0.000028	0.000672	0	0.000028
JJ	0	0	0.000340	0.000097	0
NN	0	0.000200	0.000223	0.000006	0.002337
RB	0	0	0.010446	0	0
DT	0	0	0	0.506099	0

Figure 10.6 Observation likelihoods B computed from the WSJ corpus without smoothing.

HMMS – PREDICTION (DECODING)

- ▶ $\hat{\mathbf{t}} = \operatorname{argmax}_{\mathbf{t}} P(\mathbf{w}|\mathbf{t}) P(\mathbf{t}) = \operatorname{argmax}_{\mathbf{t}} \prod_{i=1}^n P(w_i|t_i)P(t_i|t_{i-1})$
- ▶ Simple idea: for each word, take the tag that maximises $P(w_i|t_i)P(t_i|t_{i-1})$. Do it left-to-right.
- ▶ This is wrong! We are looking for $\operatorname{argmax}_{\mathbf{t}}$, not individual $\operatorname{argmax}_{t_i}$ terms.
 - ▶ This is a local classifier: error propagation
- ▶ Correct way: take **all** possible tag combinations, evaluate them, take the max (like Naïve Bayes)
 - ▶ Problem: exponential number of sequences.

THE VITERBI ALGORITHM

- ▶ Dynamic Programming to the rescue!
 - ▶ We can still proceed sequentially, as long as we careful.
- ▶ “a cat” -> a/**DT** cat/**NN**
- ▶ Best tag for “a” is easy: take $\operatorname{argmax}_t P(a|t)P(t|<s>)$
 - ▶ We can do that because first “tag” is always “<s>”
- ▶ Suppose best tag for “a” is DT. To get the tag for “cat”, we can take $\operatorname{argmax}_t P(cat|t)P(t|DT)$ but this is wrong.
- ▶ Instead, we keep track of **scores for each tag** for “a” and check **what would happen** if “a” had a different tag.

THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP					
MD					
VB					
JJ					
NN					
RB					
DT					


THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP	$P(\text{Janet} \text{NNP}) * P(\text{NNP} \langle s \rangle)$				
MD	$P(\text{Janet} \text{MD}) * P(\text{MD} \langle s \rangle)$				
VB	...				
JJ	...				
NN	...				
RB	...				
DT	...				

THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP	0.000032 * 0.2767				
MD	0 * 0.0006				
VB	...				
JJ	...				
NN	...				
RB	...				
DT	...				

THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP	8.8544e-06 				
MD	0				
VB	0				
JJ	0				
NN	0				
RB	0				
DT	0				

THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP	8.8544e-06 ●	$P(\text{will} \text{NNP}) * P(\text{NNP} t_{\text{Janet}}) * s(t_{\text{Janet}} \text{Janet})$			
MD	0	...			
VB	0	...			
JJ	0	...			
NN	0	...			
RB	0	...			
DT	0	...			

THE VITERBI ALGORITHM


	Janet	will	back	the	bill
NNP	8.8544e-06	$P(\text{will} \text{NNP}) * P(\text{NNP} t_{\text{Janet}}) * s(t_{\text{Janet}} \text{Janet})$			
MD	0	...			
VB	0				
JJ	0	...			
NN	0	...			
RB	0	...			
DT	0	...			

Calculate this for all tags,
take the max.


THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP	8.8544e-06 ●	$P(\text{NNP} t_{\text{Janet}}) \cdot s(t_{\text{Janet}} \text{Janet})$			
MD	0	...			
VB	0	...			
JJ	0	...			
NN	0	...			
RB	0	...			
DT	0	...			


THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP	8.8544e-06 	0			
MD	0	$P(\text{will} \text{MD}) * P(\text{MD} \text{t}_{\text{Janet}}) * s(\text{t}_{\text{Janet}} \text{Janet})$			
VB	0	...			
JJ	0	...			
NN	0	...			
RB	0	...			
DT	0	...			

THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP	8.8544e-06 	0			
MD	0	3.004e-8			
VB	0	...			
JJ	0	...			
NN	0	...			
RB	0	...			
DT	0	...			

THE VITERBI ALGORITHM

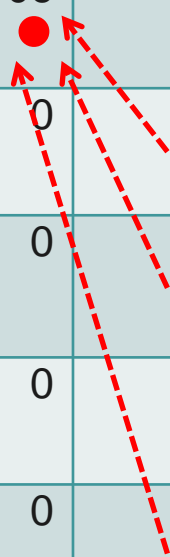
	Janet	will	back	the	bill
NNP	8.8544e-06 	0			
MD	0	3.004e-8			
VB	0	2.231e-13			
JJ	0	0			
NN	0	1.034e-10			
RB	0	0			
DT	0	0			

THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP	8.8544e-06	0			
MD	0	3.004e-8			
VB	0	2.231e-13			
JJ	0	0			
NN	0	1.034e-10			
RB	0	0			
DT	0	0			

THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP	8.8544e-06	0	0		
MD	0	3.004e-8	0		
VB	0	2.231e-13	$P(\text{back} \text{VB}) * P(\text{VB} t_{\text{will}}) * s(t_{\text{will}} \text{will})$		
JJ	0	0			
NN	0	1.034e-10			
RB	0	0			
DT	0	0			



THE VITERBI ALGORITHM

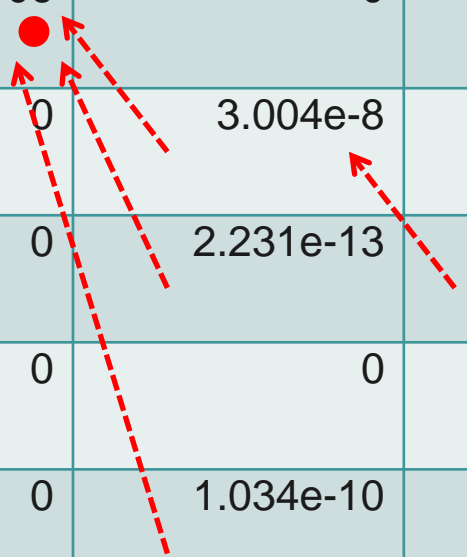
	Janet	will	back	the	bill
NNP	8.8544e-06	0	0		
MD	0	3.004e-8	0		
VB	0	2.231e-13	MD: 1.6e-11 VB: 7.5e-19 NN: 9.7e-17		
JJ	0	0			
NN	0	1.034e-10			
RB	0	0			
DT	0	0			

THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP	8.8544e-06	0	0		
MD	0	3.004e-8	0		
VB	0	2.231e-13	MD: 1.6e-11 VB: 7.5e-19 NN: 9.7e-17		
JJ	0	0			
NN	0	1.034e-10			
RB	0	0			
DT	0	0			

THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP	8.8544e-06	0	0		
MD	0	3.004e-8	0		
VB	0	2.231e-13	1.6e-11		
JJ	0	0			
NN	0	1.034e-10			
RB	0	0			
DT	0	0			



THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP	8.8544e-06	0	0		
MD	0	3.004e-8	0		
VB	0	2.231e-13	1.6e-11		
JJ	0	0	5.1e-15		
NN	0	1.034e-10	5.4e-15		
RB	0	0	5.3e-11		
DT	0	0	0		

The diagram illustrates the Viterbi algorithm for the sentence "Janet will back the bill". The table shows the probability of each word being a specific part of speech (NNP, MD, VB, JJ, NN, RB, DT). Red dashed arrows indicate the most likely path from the start to the end of the sentence, starting from the NNP cell for "Janet" and ending at the DT cell for "bill".

THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP	8.8544e-06	0	0	2.5e-17	
MD	0	3.004e-8	0	0	
VB	0	2.231e-13	1.6e-11	0	
JJ	0	0	5.1e-15	5.2e-16	
NN	0	1.034e-10	5.4e-15	5.9e-18	
RB	0	0	5.3e-11	0	
DT	0	0	0	1.8e-12	

The diagram shows red dashed arrows indicating the backtracking path for the Viterbi algorithm. The path starts at the NNP state for 'Janet' (8.8544e-06), which is marked with a red dot. Arrows point from this state to the MD state for 'Janet' (0), then to the VB state for 'Janet' (0), then to the JJ state for 'Janet' (0), then to the NN state for 'Janet' (0), then to the RB state for 'Janet' (0), and finally to the DT state for 'Janet' (0). This path represents the most likely sequence of states for the given words.

THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP	8.8544e-06	0	0	2.5e-17	0
MD	0	3.004e-8	0	0	0
VB	0	2.231e-13	1.6e-11	0	1.0e-20
JJ	0	0	5.1e-15	5.2e-16	0
NN	0	1.034e-10	5.4e-15	5.9e-18	2.0e-15
RB	0	0	5.3e-11	0	0
DT	0	0	0	1.8e-12	0

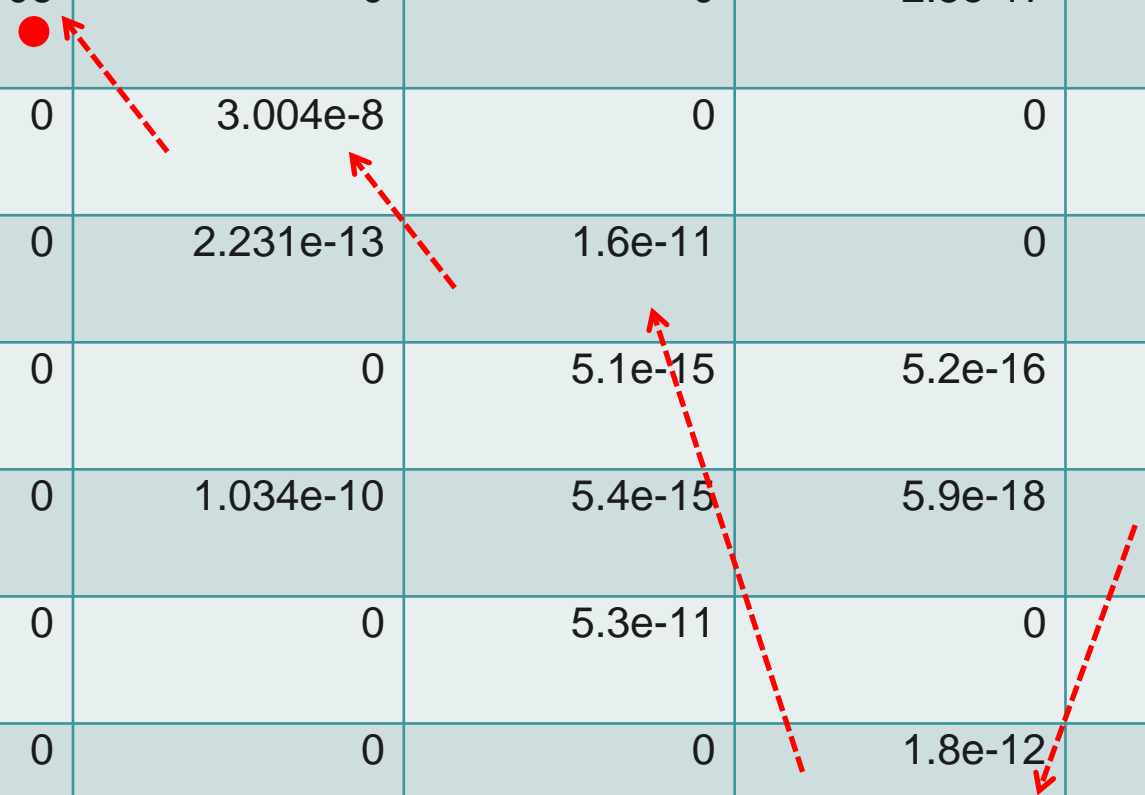
The diagram shows red dashed arrows indicating the backpointers for the sequence "Janet will back the bill". A red dot is placed at the NNP tag for "Janet". The arrows trace the path of the most likely sequence, starting from the NNP tag for "Janet" and moving through the MD tag for "will", the VB tag for "back", the JJ tag for "the", and finally the DT tag for "bill".

THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP	8.8544e-06	0	0	2.5e-17	0
MD	0	3.004e-8	0	0	0
VB	0	2.231e-13	1.6e-11	0	1.0e-20
JJ	0	0	5.1e-15	5.2e-16	0
NN	0	1.034e-10	5.4e-15	5.9e-18	2.0e-15
RB	0	0	5.3e-11	0	0
DT	0	0	0	1.8e-12	0

THE VITERBI ALGORITHM

	Janet	will	back	the	bill
NNP	8.8544e-06	0	0	2.5e-17	0
MD	0	3.004e-8	0	0	0
VB	0	2.231e-13	1.6e-11	0	1.0e-20
JJ	0	0	5.1e-15	5.2e-16	0
NN	0	1.034e-10	5.4e-15	5.9e-18	2.0e-15
RB	0	0	5.3e-11	0	0
DT	0	0	0	1.8e-12	0



THE VITERBI ALGORITHM

	Janet/ NNP	will/ MD	back/ VB	the/ DT	bill/ NN
NNP	8.8544e-06	0	0	2.5e-17	0
MD	0	3.004e-8	0	0	0
VB	0	2.231e-13	1.6e-11	0	1.0e-20
JJ	0	0	5.1e-15	5.2e-16	0
NN	0	1.034e-10	5.4e-15	5.9e-18	2.0e-15
RB	0	0	5.3e-11	0	0
DT	0	0	0	1.8e-12	0

The diagram illustrates the Viterbi algorithm's backtracking process. Red dashed arrows trace the path of maximum probability from the final state (NN, 2.0e-15) back to the initial state (NNP, 8.8544e-06). The path consists of the following states: (NN, 2.0e-15) → (DT, 1.8e-12) → (VB, 1.6e-11) → (MD, 3.004e-8) → (NNP, 8.8544e-06). A red dot is placed on the final state (NN, 2.0e-15) to indicate the end of the path.

THE VITERBI ALGORITHM

- ▶ Complexity: $O(T^2N)$, where T is the size of the tagset and N is the length of the sequence.
 - ▶ $T * N$ matrix, each cell performs T operations.
- ▶ Why does it work?
 - ▶ Because of the **independence assumptions** that decompose the problem (specifically, the Markov property). Without these, we cannot apply DP.

VITERBI PSEUDOCODE

```
alpha = np.zeros(M, T)
for t in range(T):
    alpha[1, t] = pi[t] * O[w[1], t]

for i in range(2, M):
    for t_i in range(T):
        for t_last in range(T):      # t_last means t_{i-1}
            s = alpha[i-1, t_last] * A[t_last, t_i] * O[w[i], t_i]
            if s > alpha[i, t_i]:
                alpha[i, t_i] = s
                back[i, t_i] = t_last

best = np.max(alpha[M-1, :])
return backtrace(best, back)
```

- ▶ Good practice: work with **log** probabilities to prevent underflow (multiplications become sums)
- ▶ Vectorisation (use matrix-vector operations)

HMMS IN PRACTICE

- ▶ We saw HMM taggers based on **bigrams**. State-of-the-art use tag **trigrams**.
 - ▶ $P(\mathbf{t}) = \prod_{i=1}^n P(t_i | t_{i-1}, t_{i-2})$ Viterbi now $O(T^3N)$
- ▶ Need to deal with sparsity: some tag trigram sequences might not be present in training data
 - ▶ Backoff: $P(t_i | t_{i-1}, t_{i-2}) = \lambda_3 \hat{P}(t_i | t_{i-1}, t_{i-2}) + \lambda_2 \hat{P}(t_i | t_{i-1}) + \lambda_1 \hat{P}(t_i)$
 - ▶ $\lambda_1 + \lambda_2 + \lambda_3 = 1$
 - ▶ Can learn the weights using **deleted interpolation**.
- ▶ With additional features, reach 96.7% accuracy on Penn Treebank (Brants, 2000)

OTHER VARIANT TAGGERS

- HMM is **generative**, $P(t, w)$, ‘creates’ the input
 - allows for unsupervised HMMs: learn model without any tagged data!
- **Discriminative** models also popular, modelling $P(t \mid w)$ directly
 - supports richer feature set, generally better accuracy when trained over large supervised datasets
 - E.g., Maximum Entropy Markov Model (MEMM), Conditional random field (CRF), Connectionist Temporal Classification (CTC)
 - Most *deep learning* models of sequences are discriminative (e.g., encoder-decoders for translation), similar to an MEMM

HMMS IN NLP

- HMMS are highly effective for part-of-speech tagging
 - trigram HMM gets 96.7% accuracy (TnT)
 - related models are state of the art
 - MEMMs 97.3%
 - CRFs 97.6%
 - *English Penn Treebank* tagging accuracy
[https://aclweb.org/aclwiki/index.php?title=POS Tagging \(State of the art\)](https://aclweb.org/aclwiki/index.php?title=POS_Tagging_(State_of_the_art))
- Other sequence labelling tasks
 - named entity recognition, shallow parsing, alignment ...
 - In other fields: DNA, protein sequences, image lattices...

A FINAL WORD

- ▶ HMMs are a simple, yet effective way to perform sequence labelling.
- ▶ Can still be competitive, and fast. Natural baseline for other sequence labelling tasks.
- ▶ Main drawback: not very flexible in terms of feature representation, compared to MEMMs and CRFs.
- ▶ Next lecture: unsupervised HMMs.

READINGS

- ▶ Hidden Markov models
 - ▶ JM3 Ch 9.1, 9.2, 9.4, 10.4
- ▶ [Optional] Rabiner's HMM tutorial, for more details
 - ▶ <http://tinyurl.com/2hqaf8>
- ▶ [Optional] Contemporary sequence tagging methods
 - ▶ JM3 Ch 10.5-10.6 - MEMMs
 - ▶ Lafferty et al, Conditional random fields: Probabilistic models for segmenting and labeling sequence data (2001), ICML - CRFs