

lots of |

lots of love

lots of fish

lots of discharge

lots of lollies

Press Enter to search.

COMP90042 LECTURE 14

NEURAL LANGUAGE MODELS

LANGUAGE MODELS

- ▶ Assign a probability to a sequence of words
- ▶ Framed as “sliding a window” over the sentence, predicting each word from finite context to left

E.g., $n = 3$, a trigram model

$$P(w_1, w_2, \dots, w_m) = \prod_{i=1}^m P(w_i | w_{i-2} w_{i-1})$$

- ▶ Training (estimation) from frequency counts
 - ▶ Difficulty with rare events \rightarrow smoothing
- ▶ But are there better solutions?
Neural networks (deep learning) a popular alternative.

OUTLINE

- ▶ Neural network fundamentals
- ▶ “Feed-forward” neural language models
- ▶ Recurrent neural language models

LMS AS CLASSIFIERS

LMS can be considered simple classifiers, e.g. trigram model

$$P(w_m | w_{m-2} = \text{"cow"}, w_{m-1} = \text{"eats"})$$



classifies the likely next word in a sequence.

Has a parameter for every combination of

$$w_{m-2}, w_{m-1}, w_m$$

Can think of this as a specific type of classifier — one with a very simple parameterisation.


BIGRAM LM IN PICTURES



FROM BIGRAMS TO TRIGRAMS


► Problems

- too many parameters (V^3) to be learned from data (growing with LM order)
- parameters completely separate for different trigrams
 - even when sharing words, e.g., (cow, eats) vs (farmer, eats) [*motivating smoothing*]
 - doesn't recognise similar words e.g., cow vs sheep, eats vs chews




CAN WE USE WORD EMBEDDINGS?

- ▶ How about embedding the words?
- ▶ Fewer parameters:
 - ▶ embedding matrix (or two matrices), of size $V \times d$
 - ▶ weights now $d \times d$
- ▶ $d \ll V$ is a parameter, typically in $[100, 1000]$
- ▶ Model is the *log-bilinear LM* (closely related to *word2vec*)



FEED FORWARD NEURAL NET LMS

- ▶ Neural networks more general approach, based on same principle
 - ▶ embed input context words
 - ▶ transform in “hidden” space
 - ▶ un-embed to prob over full vocab
- ▶ Neural network used to define transformations
 - ▶ e.g., feed forward LM (FFLM)



INTERIM SUMMARY

- ▶ Ngram LMs
 - ▶ cheap to train (just compute counts)
 - ▶ but too many parameters, problems with sparsity and scaling to larger contexts
 - ▶ don't adequately capture properties of words (grammatical and semantic similarity), e.g., film vs movie
- ▶ NNLMs more robust
 - ▶ force words through low-dimensional embeddings
 - ▶ automatically capture word properties, leading to more robust estimates

NEURAL NETWORKS


“Deep” neural networks provide mechanism for learning richer models.

Based on **vector** *embeddings* and compositional functions over these vectors.

- ▶ *Word embeddings capture* grammatical and semantic similarity “cows” ~ “sheep”, “eats” ~ “chews” etc.
- ▶ Vector composition can allow for combinations of features to be learned (e.g., humans consume meat)
- ▶ Limit size of vector representation to keep model capacity under control.

COMPONENTS OF NN CLASSIFIER

- ▶ NN = Neural Network
 - ▶ a.k.a. artificial NN, deep learning, multilayer perceptron
- ▶ Composed of simple functions of vector-valued inputs



NN UNITS

- ▶ Each “unit” is a function

- ▶ given input x , computes real-value (scalar) h

$$h = \tanh \left(\sum_j w_j x_j + b \right)$$

- ▶ scales input (with weights, w) and adds offset (bias, b)
 - ▶ applies non-linear function, such as logistic sigmoid, hyperbolic sigmoid (\tanh), or rectified linear unit

NEURAL NETWORK COMPONENTS

- ▶ Typically have several hidden units, i.e.,

$$h_i = \tanh \left(\sum_j w_{ij} x_j + b_i \right)$$



- ▶ each with own weights (w_i) and bias term (b_i)
- ▶ can be expressed using matrix & vector operators

$$\vec{h} = \tanh \left(W\vec{x} + \vec{b} \right)$$

- ▶ where W is a matrix comprising the unit weight vectors, and b is a vector of all the bias terms
- ▶ and *tanh* applied element-wise to a vector
- ▶ Very efficient and simple implementation


ANN IN PICTURES

- ▶ Pictorial representation of a single unit, for computing y from x
- ▶ Typical networks have several units, and additional layers
- ▶ E.g., output layer, for classification target



BEHIND THE NON-LINEARITY


- ▶ Non-linearity lends expressive power beyond logistic regression (linear ANN is otherwise equivalent)



- ▶ Non-linearity allows for complex functions to be learned
 - ▶ e.g., logical operations: AND, OR, NOT etc.
 - ▶ single hidden layer ANN is *universal approximator*: can represent any function with sufficiently large hidden state


EXAMPLE NETWORKS

- ▶ What do these units do?
 - ▶ using the “step” activation function
 - ▶ consider binary values for x_1 and x_2



COUPLING THE OUTPUT LAYER

- ▶ To make this into a classifier, need to produce a classification output
 - ▶ e.g., vector of probabilities for the next word
- ▶ Add another layer, which takes h as input, and maps to classification target (e.g., V)
 - ▶ to ensure probabilities, apply “softmax” transform
 - ▶ softmax applies exponential, and normalises, i.e., applied to vector \mathbf{v}



$$\vec{h} = \tanh \left(W \vec{x} + \vec{b} \right)$$

$$\vec{y} = \text{softmax} \left(U \vec{h} \right)$$

$$\left[\frac{\exp(v_1)}{\sum_i \exp(v_i)}, \frac{\exp(v_2)}{\sum_i \exp(v_i)}, \dots, \frac{\exp(v_m)}{\sum_i \exp(v_i)} \right]$$

DEEP STRUCTURES


- ▶ Can stack several hidden layers; e.g.,
 1. map from 1-hot words, w , to word embeddings, e (lookup)
 2. transform e to hidden state h_1 (with non-linearity)
 3. transform h_1 to hidden state h_2 (with non-linearity)
 4. ... *repeat* ...
 5. transform h_n , to output classification space y (with softmax)
- ▶ Each layer typically fully-connected to next lower layer, i.e., each unit is connected to all input elements
- ▶ Depth allows for more complex functions, e.g., longer logical expressions

LEARNING WITH SGD

- ▶ How to learn the parameters from data?
 - ▶ parameters = sets of weights, bias, embeddings
- ▶ Consider how well the model “fits” the training data, in terms of the probability it assigns to the correct output
 - ▶ *e.g.*, $\prod_{i=1}^m P(w_i | w_{i-2} w_{i-1})$ for sequence of m words
 - ▶ want to *maximise* this probability, equivalently *minimise* its negative log
 - ▶ $-\log P$ is known as the “loss” function
- ▶ Trained using gradient based methods, much like logistic regression (tools like *tensorflow*, *theano*, *torch* etc use autodiff to compute gradients automatically)

FFNNLM

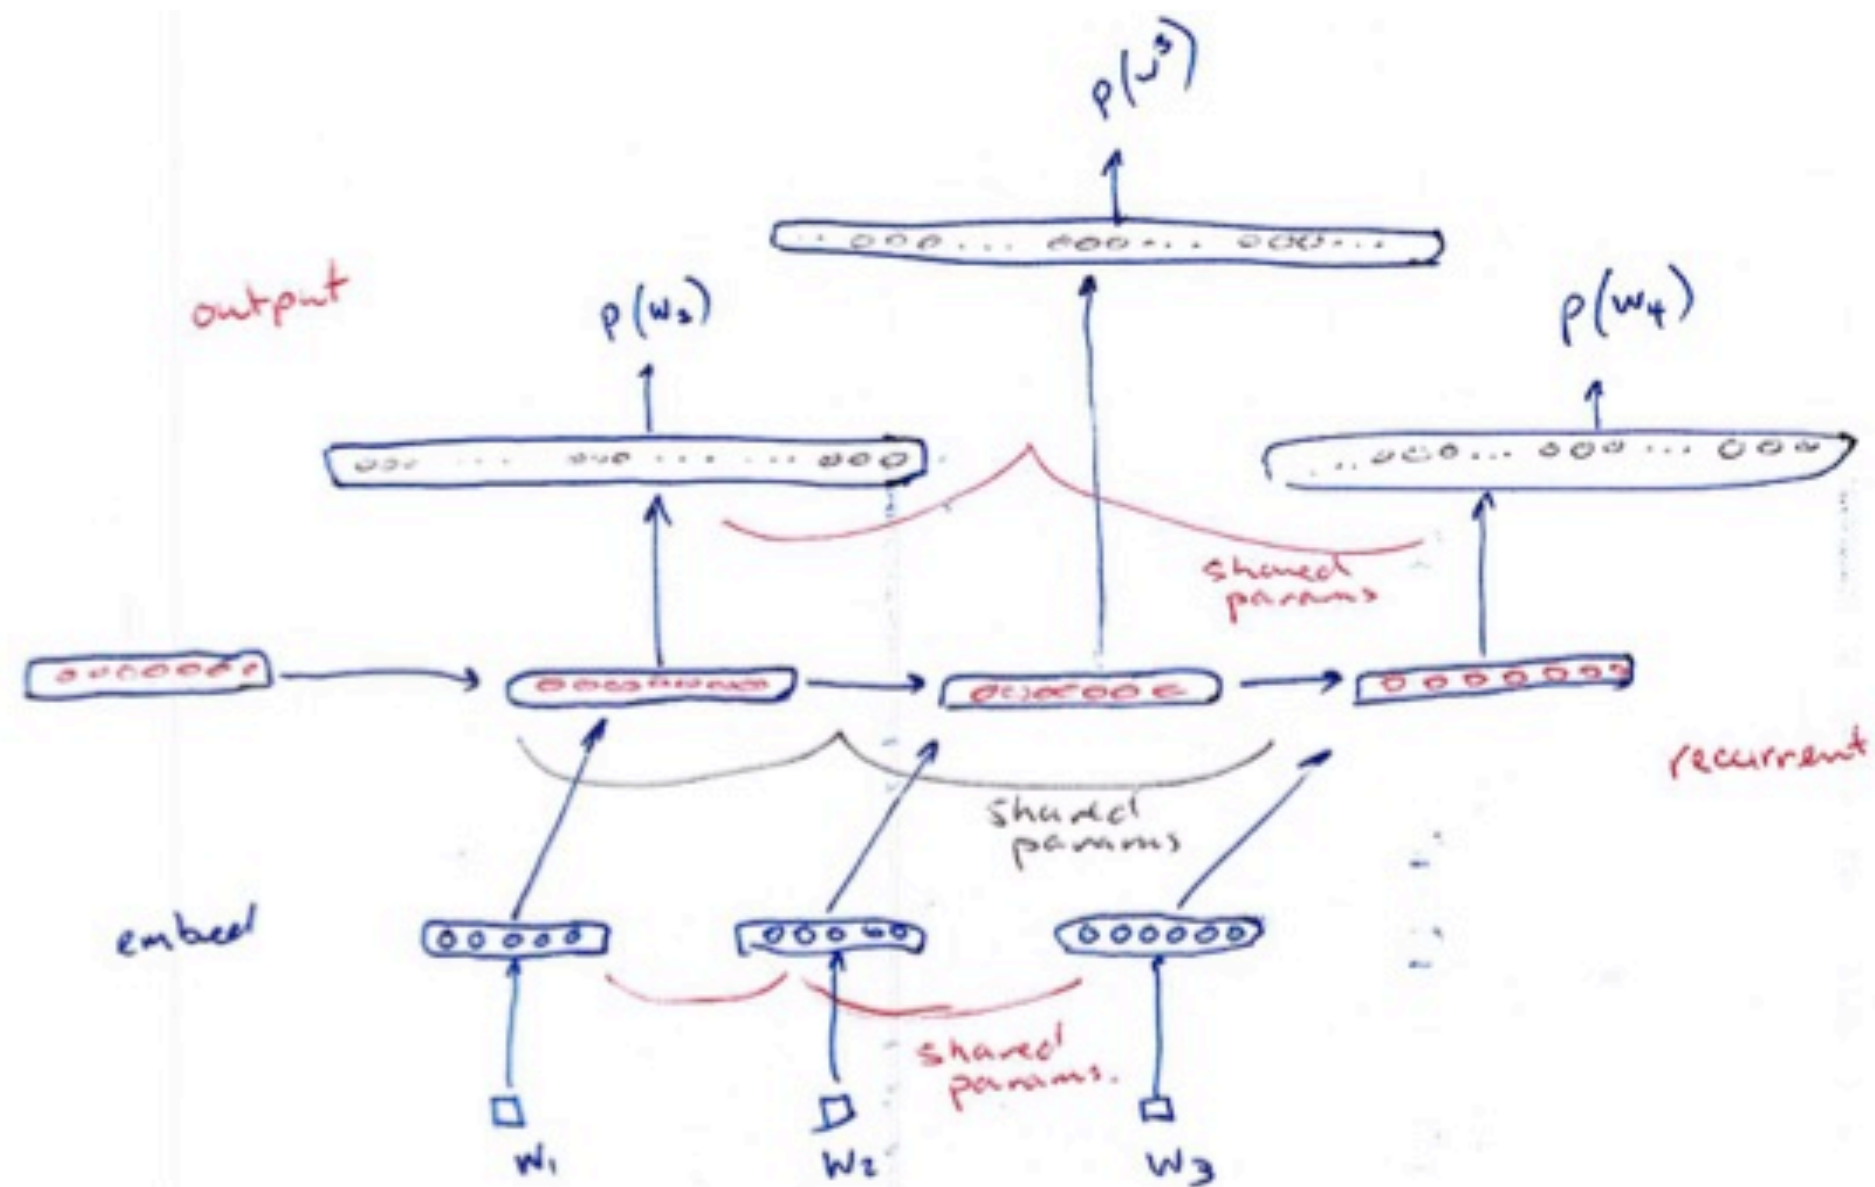
► Application to language modelling



Bengio et al, 2003


RECURRENT>NNLMS

- What if we structure the network differently, e.g., according to sequence with Recurrent Neural Networks (RNNs)



RECURRENT>NNLMS

- ▶ Start with
 - ▶ initial hidden state \mathbf{h}_0
- ▶ For each word, w_i , in order $i=1..m$
 - ▶ embed word to produce vector, \mathbf{e}_i
 - ▶ compute hidden $\mathbf{h}_i = \tanh(W \mathbf{e}_i + V \mathbf{h}_{i-1} + \mathbf{b})$
 - ▶ compute output $P(\mathbf{w}_{i+1}) = \text{softmax}(U \mathbf{h}_i + \mathbf{c})$
- ▶ Train such to minimise $\sum_i -\log P(\mathbf{w}_i)$
 - ▶ to learn parameters $W, V, U, \mathbf{b}, \mathbf{c}, \mathbf{h}_0$



RNNS

- ▶ Can result in very “deep” networks, difficult to train due to gradient explosion or vanishing
 - ▶ variant RNNs designed to behave better, e.g., GRU, LSTM
- ▶ RNNs used widely as sentence encodings
 - ▶ RNN processes sentence, word at a time, use final state as fixed dimensional representation of sentence (of any length)
 - ▶ Can also run another RNN over reversed sentence, and concatenate both final representations
 - ▶ Used in translation, summarisation, generation, text classification, and more

FINAL WORDS

- ▶ NNet models
 - ▶ Robust to word variation, typos, etc
 - ▶ Excellent generalization, especially RNNs
 - ▶ Flexible — forms the basis for many other models (translation, summarization, generation, tagging, etc)
- ▶ Cons
 - ▶ Much slower than counts... but hardware acceleration
 - ▶ Need to limit vocabulary, not so good with rare words
 - ▶ Not as good at memorizing fixed sequences
 - ▶ Data hungry, not so good on tiny data sets

REQUIRED READING

- ▶ J&M3 Ch. 8
- ▶ Neubig 2017, “Neural Machine Translation and Sequence-to-sequence Models: A Tutorial”, Sections 5 & 6