

Query Completion / Expansion

COMP90042 LECTURE 18, THE UNIVERSITY OF MELBOURNE

by

Matthias Petri

What is a query?

What is a query?

1. Obviously the stuff I type into the search box!
2. Most likely not the query that gets handed over to the search index.
3. Why not?

Query Completion

Query Completion

why is my|



why is my **internet so slow**

why is my **computer so slow**

why is my **hair falling out**

why is my **period late**

why is my **printer offline**

why is my **poop green**

why is my **mac so slow**

why is my **laptop so slow**

why is my **ipad so slow**

why is my **phone so hot**

why is my **dog eating grass**

What is a Query Completion?

Goals:

1. Assist users to formulate search requests.
2. Reduce number of keystrokes required to enter query.
3. Help with spelling query terms.
4. Guide user towards what a good query might be.
5. Cache results! Reduce server load.

Strategy:

1. Generate list of completions based on partial query.
2. Refine suggestions as more keys are pressed.
3. Stop once users selects candidate or completion fails.
4. Why not a Language Model? Might not return results!

High Level Algorithm

Given a query pattern P ,

1. Retrieve set of **candidates** “matching” P from set S of possible target queries.
2. Rank candidates by frequency.
3. Possibly re-rank highest ranked candidates with more complex ranking measure (e.g. personalized)
4. Return the top- K highest ranking candidates as suggestions.

Completion Targets

Where does the set S of possible completions come from?

1. Most popular queries (websearch)
2. Items listed on website (ecommerce)
3. Past queries by the user (email search)

Properties:

1. Static (e.g. completion for “twi”)
2. Dynamic (e.g. time-sensitive, “world cup”)
3. Massive or small (email search vs websearch)

Completion Types / Completion

Given a partial user query P , how is the initial candidate set retrieved?

Modes:

1. Prefix match.
2. Substring match.
3. Multi-term prefix match.
4. Relaxed match.

Example: Target “FIFA world cup 2018”:

P	Mode 1	Mode 2	Mode 3	Mode 4
FIFA wo	x	x	x	x
orl		x		
FI wor			x	x
FIFO world cu				x

Prefix Completion

Problem:

Given a query prefix P , retrieve the top- K most popular completions.

Data:

Static query log consisting of all queries received by the search index.

Requirements:

1. Fast retrieval time required. What is fast?
2. Space efficient index.

Prefix match - Trie+RMQ based Index

Step 1: Preprocess data by sorting query log in lexicographical order and counting frequency of unique queries:

Before	After
bunnings	< <i>bunnings</i> , 47 >
bachelor in paradise	< <i>big w</i> , 5 >
bbc news	< <i>bbc news</i> , 12 >
bunnings	< <i>bachelor in paradise</i> , 2 >
big w	
bbc news	
big w	

Prefix match - Trie+RMQ based Index

Step 2: Insert all unique queries and their frequencies into a trie (also called a prefix tree).

What is a trie?

- A tree representing a set of strings.
- Edges of the tree are labeled.
- Children of nodes are ordered.
- Root to node path represents prefix of all strings in the subtree starting at that node.

Prefix match - Trie Example

Set of strings:

- nba
- news
- nab
- ngv
- netflix
- netbank
- network
- netball
- netbeans

<https://www.cs.usfca.edu/~galles/visualization/Trie.html>

Prefix match - Trie+RMQ based Index

Prefix search using a trie

Insert queries into trie. For a pattern P , find node in trie representing the subtree prefixed by P in $O(|P|)$ time.


Observation:

The subtree prefixed by P corresponds to a continuous range.

Prefix match - Trie+RMQ based Index

Idea:

Store array with frequencies corresponding to each query. Subtree corresponds to range in frequency array. Find the top-K highest numbers in that range.



Range Maximum Queries

Task:

Given an array A of n numbers, and a range $[l, r]$ of size m , find the **positions** of the K largest numbers in $A[l, r]$.

Simple algorithm:

1. Copy $A[l, r]$ into an array B in $O(m)$ time.
2. Sort B in $O(m \log m)$ time.
3. Return positions of largest numbers in $A[l, r]$.

Problem:

- Runtime also depends on the size of the range m and requires $O(m)$ extra space.
- m can be large. We require low millisecond response times.

Range Maximum Queries - Index

Finding the Maximum in a Range in $O(1)$ time:

- Array A is size n .
- There are $O(n^2)$ different ranges $A[i, j]$
- For each range precompute the position of the maximum. Uses $O(n^2)$ space.

Extension to K largest numbers:

1. Find position p of largest element on $A[i, j]$.
2. Recurse to $A[i, p - 1]$ and $A[p + 1, j]$.
3. Keep going until you have the K largest elements.
4. Runtime $O(K \log m)$.

RMQ Index- Reduce space

Simple space reduction:

- Instead of precomputing all $O(n^2)$ ranges $A[i, j]$, for each position $A[i]$, precompute only $\log n$ ranges of increasing size: $A[i, i + 1], A[i, i + 2], A[i, i + 4], A[i, i + 8]$.
- Any range $A[l, r]$ can be decomposed into two ranges $A[l, Y]$ and $A[Z, r]$ where $Y = l + 2^x$ and $Z = r - 2^y$ such that $Z \geq l$, $Y \leq r$ and, $A[l, Y], A[Z, r]$ overlap. Then,
 $RMQ(A[i, j]) = \max(RMQ(A[l, Y]), RMQ(A[Z, r]))$
- Total space cost $O(n \log n)$.

Prefix Completion - In Practice

- Space efficient (compressed) Trie+RMQ representations used (more complex)
- RMQ+Trie requires roughly 10 bytes per string (roughly the size of gzip).
- 1 billion unique strings require an index of size 10GB RAM.
- Can answer top-10 queries in less than 10 microseconds.

Query Expansion

Query Expansion - What is it?

- User and documents may refer to a concept using different words (poison \leftrightarrow toxin, danger \leftrightarrow hazard, postings list \leftrightarrow inverted list)
- Vocabulary mismatch can have impact on recall
- Users often attempt to fix this problem manually (query reformulation)
- Adding these synonyms should improve query performance (query expansion)

Global Query Expansion

- Retrieve synonyms from thesaurus or WordNet (medical domain)
- Spell correction (important → important)
- Word2Vec (what words are close to the query words?)

PubMed [Create RSS](#) [Create alert](#) [Advanced](#) [Help](#)

Format: Summary ▾ Sort by: Most Recent ▾ Per page: 20 ▾ Send to ▾ Filters: [Manage Filters](#)

Search results
Items: 1 to 20 of 119

<< First < Prev Page 1 of 6 Next > Last >>

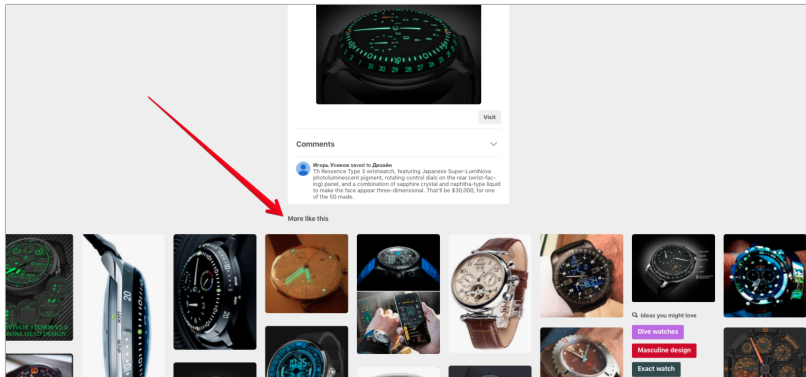
☐ [Comprehensive systematic review summary: Disease-modifying therapies for adults with multiple sclerosis: Report of the Guideline Development, Dissemination, and Implementation Subcommittee of the American Academy of Neurology.](#)
Rae-Grant A, Day GS, Marrie RA, Rabinstein A, Cree BAC, Gronseth GS, Haboubi M, Halper J, Hosey JP, Jones DE, Lisak R, Pelletier D, Potrebic S, Sitcov C, Sommers R, Stachowiak J, Getchius TSD, Merillat SA, Pringsheim T.
Neurology. 2018 Apr 24;90(17):789-800. doi: [10.1212/WNL.0000000000005345](#).
PMID: [29686117](#)

Search details

```
("ocrelizumab"[Supplementary Concept OR "ocrelizumab"[All Fields]) AND ("multiple sclerosis"[MeSH Terms] OR ("multiple"[All Fields] AND "sclerosis"[All Fields]) OR "multiple sclerosis"[All Fields])
```

User relevance feedback

Relevance Feedback. User provides feedback to the search engine by indicating which results are relevant



Pseudorelevance feedback

- Take top- K results of original query
- Determine important/informative terms/topics (topic modelling!) shared by those documents
- Expand query by those terms
- No explicit user feedback needed (also called blind relevance feedback)

Example

- Original query: what is a prime factors
- Expanded query: what is a prime factors integer number composite common divisor

Indirect relevance feedback

- For a query look at what users click on in the result page
- Use clicks as signal of relevance
- Learning-2-Rank uses neural models to rerank result pages (later this semester)

Query Expansion - Summary

- Helps with vocabulary mismatch
- Can improve recall
- Global expansion
- User, pseudo or indirect relevance feedback

Further Reading

Reading:

- Manning, Christopher D; Raghavan, Prabhakar; Schütze, Hinrich; Introduction to information retrieval, Cambridge University Press 2008. (Chapter 9)

Additional References:

- Unni Krishnan, Alistair Moffat, Justin Zobel: A Taxonomy of Query Auto Completion Modes. ADCS 2017: 6:1-6:8
- Amati, Giambattista (2003) Probability models for information retrieval based on divergence from randomness. PhD thesis.