

COMP90042 Lecture 19

Index compression and efficient query processing

Matthias Petri

Wed 10/5/17

What we'll learn today

- ▶ Efficient storage and access of **Posting lists**
- ▶ Efficient inverted index query processing

Index compression

Benefits of index compression:

- ▶ Reduce storage requirements
- ▶ Keep larger parts of the index in memory
- ▶ Faster query processing

Example: A state-of-the-art inverted index of 25 million websites (426GB raw data) requires only 5GB (1.2%) and can answer queries in ≈ 10 milliseconds.

Practice: Companies such as Facebook keep their index in main memory. Efficient, compact representations can save lots of money!

Compression Principles

- ▶ Compressibility is bounded by the information content of a data set
- ▶ Information content of a text T is characterized by its *Entropy* H :

$$H(T) = - \sum_{s \in \Sigma} \frac{f_s}{n} \log_2 \frac{f_s}{n}$$

where f_s is the frequency of symbol s in T and n is the length of T

- ▶ For example, $H(\text{abracadabra}) = 2.040373$ bits with $n = 11, f_a = 5, f_b = 2, f_c = 1, f_d = 1, f_r = 2$
- ▶ Intuition: Spend less bits on items that occur often

Posting list Compression - Terminology

Inverted index terminology : Terms, Documents, Postings Lists, Postings

Example

the	25	26	29	...	12345	12347	12349
house	5123	5234	5454	5591	...		
aeronaut	251235	251239	251273				

Posting list Compression - Requirements

- ▶ Minimize storage costs
- ▶ Fast access

Sample collection

Documents N	24,622,347
Terms	35,636,425
Postings	5,742,630,292
Uncompressed Storage Cost ¹	16.71 GB

¹Only document identifier, no frequencies or positions

Posting list Compression - Concepts

- ▶ Postings list corresponds to an increasing sequence of integers
- ▶ Each integer can be in $[1, N]$ requiring $\lceil \log_2(N) \rceil$ bits.
Example: $N = 24,622,347$. $\log_2(N) = 24.55346 = 25$ bits
- ▶ Idea: Gaps between two adjacent integers can be much smaller

the	ids:	25	26	29	...	12345	12347	12349
	gaps:	25	1	3	...	1	2	2
house	ids:	5123	5234	5454	5591	...		
	gaps:	5123	1	220	137	...		
aeronaut	ids:	251235	251239	251273				
	gaps:	251235	4	34				

Variable Byte Compression

Idea

Use variable number of bytes to represent integers. Each byte contains 7 bits “payload” and one continuation bit.

Examples

Number		Bits	Encoding	
824	1100111000		00111000	10000110
5	101		10000101	

Variable Byte Compression - Storage Cost

Storage Cost

Range Start	Range Stop	Power of 2 Range		Number of Bytes
0	127	0	$2^7 - 1$	1
128	16,383	2^7	$2^{14} - 1$	2
16,384	2,097,151	2^{14}	$2^{21} - 1$	3
2,097,152	268,435,455	2^{21}	$2^{28} - 1$	4
268,435,456	34,359,738,368	2^{28}	$2^{35} - 1$	5
\vdots	\vdots	\vdots	\vdots	\vdots

Variable Byte Compression - Example

Compress number 512312 or 1111101000100111000 in binary.

- ▶ How many bytes?

Variable Byte Compression - Example

Compress number 512312 or 1111101000100111000 in binary.

- ▶ How many bytes? 11111|0100010|0111000. 3 bytes!

Variable Byte Compression - Example

Compress number 512312 or 1111101000100111000 in binary.

- ▶ How many bytes? 11111|0100010|0111000. 3 bytes!
- ▶ How do we compress the number?

Variable Byte Compression - Example

Compress number 512312 or 1111101000100111000 in binary.

- ▶ How many bytes? 11111|0100010|0111000. 3 bytes!
- ▶ How do we compress the number?
 1. Extract the lowest 7 bits.

Variable Byte Compression - Example

Compress number 512312 or 1111101000100111000 in binary.

- ▶ How many bytes? 11111|0100010|0111000. 3 bytes!
- ▶ How do we compress the number?

1. Extract the lowest 7 bits.

$$512312 \bmod 128 = 56 = 0111000$$

Variable Byte Compression - Example

Compress number 512312 or 1111101000100111000 in binary.

- ▶ How many bytes? 11111|0100010|0111000. 3 bytes!
- ▶ How do we compress the number?
 1. Extract the lowest 7 bits.
 $512312 \bmod 128 = 56 = 0111000$
 2. Discard lowest 7 bits.

Variable Byte Compression - Example

Compress number 512312 or 1111101000100111000 in binary.

- ▶ How many bytes? 11111|0100010|0111000. 3 bytes!
- ▶ How do we compress the number?
 1. Extract the lowest 7 bits.
 $512312 \bmod 128 = 56 = 0111000$
 2. Discard lowest 7 bits.
 $512312 \div 128 = 4002$ (or $512312 >> 7$)

Variable Byte Compression - Example

Compress number 512312 or 1111101000100111000 in binary.

- ▶ How many bytes? 11111|0100010|0111000. 3 bytes!
- ▶ How do we compress the number?
 1. Extract the lowest 7 bits.
 $512312 \bmod 128 = 56 = 0111000$
 2. Discard lowest 7 bits.
 $512312 \div 128 = 4002$ (or $512312 >> 7$)
 3. Extract the lowest 7 bits.

Variable Byte Compression - Example

Compress number 512312 or 1111101000100111000 in binary.

- ▶ How many bytes? 11111|0100010|0111000. 3 bytes!
- ▶ How do we compress the number?
 1. Extract the lowest 7 bits.
 $512312 \bmod 128 = 56 = 0111000$
 2. Discard lowest 7 bits.
 $512312 \div 128 = 4002$ (or $512312 >> 7$)
 3. Extract the lowest 7 bits.
 $4002 \bmod 128 = 34 = 0100010$

Variable Byte Compression - Example

Compress number 512312 or 1111101000100111000 in binary.

► How many bytes? 11111|0100010|0111000. 3 bytes!

► How do we compress the number?

1. Extract the lowest 7 bits.

$$512312 \bmod 128 = 56 = 0111000$$

2. Discard lowest 7 bits.

$$512312 \div 128 = 4002 \text{ (or } 512312 \gg 7)$$

3. Extract the lowest 7 bits.

$$4002 \bmod 128 = 34 = 0100010$$

4. Discard lowest 7 bits.

$$4002 \div 128 = 31 \text{ (or } 4002 \gg 7)$$

Variable Byte Compression - Example

Compress number 512312 or 1111101000100111000 in binary.

- ▶ How many bytes? 11111|0100010|0111000. 3 bytes!
- ▶ How do we compress the number?
 1. Extract the lowest 7 bits.
 $512312 \bmod 128 = 56 = 0111000$
 2. Discard lowest 7 bits.
 $512312 \div 128 = 4002$ (or $512312 >> 7$)
 3. Extract the lowest 7 bits.
 $4002 \bmod 128 = 34 = 0100010$
 4. Discard lowest 7 bits.
 $4002 \div 128 = 31$ (or $4002 >> 7$)
 5. Number smaller than 128. Write in lowest 7 bits and set top bit to 1. $31 = 11111$ So we write 10011111 which is $31 + 128 = 159$

Variable Byte Compression - Example

Compress number 512312 or 1111101000100111000 in binary.

- ▶ How many bytes? 11111|0100010|0111000. 3 bytes!
- ▶ How do we compress the number?
 1. Extract the lowest 7 bits.
 $512312 \bmod 128 = 56 = 0111000$
 2. Discard lowest 7 bits.
 $512312 \div 128 = 4002$ (or $512312 \gg 7$)
 3. Extract the lowest 7 bits.
 $4002 \bmod 128 = 34 = 0100010$
 4. Discard lowest 7 bits.
 $4002 \div 128 = 31$ (or $4002 \gg 7$)
 5. Number smaller than 128. Write in lowest 7 bits and set top bit to 1. $31 = 11111$ So we write 10011111 which is $31 + 128 = 159$

Variable Byte Compression - Algorithm

Encoding

```
1: function ENCODE( $x$ )  
2:   while  $x \geq 128$  do  
3:     WRITE( $x \bmod 128$ )  
4:      $x = x \div 128$   
5:   end while  
6:   WRITE( $x + 128$ )  
7: end function
```

Decoding

```
1: function DECODE(bytes)  
2:    $x = 0$   
3:    $y = \text{READBYTE}(\text{bytes})$   
4:   while  $y < 128$  do  
5:      $x = 128 \times x + y$   
6:      $y = \text{READBYTE}(\text{bytes})$   
7:   end while  
8:    $x = 128 \times x + (y - 128)$   
9:   return  $x$   
10: end function
```

OptPForDelta Compression


Idea

Group k gaps and encode using fixed number of bits. Encode numbers $> 2^b$ separately as an *exception*. Pick b “*optimally*” for each block so there are $\approx 10\%$ exceptions.

Example $k = 8$

[1 4 7 2 4 5 123 6] [1 4 5 232 523 7 2 3] [3 4 755 15 12 1 8 4]
 b=3 b=8 b=4

Encode [1 4 7 2 4 5 123 6] as:



Postings List Decompression Speeds/Space Usage

Algorithm	Space [bits/posting]	Speed [Million Integers/sec]
Uncompressed	32	≈ 5400
Variable Byte	8.7	≈ 680
OptPForDelta	4.7	≈ 710
Simple-8b	4.8	≈ 780
SIMD-BP128	11	≈ 2300
\vdots	\vdots	

Postings List Compression - Optimizations

- ▶ Commonly lists are split into blocks of 128 integers
- ▶ Choose optimal compression for each block
- ▶ Often, long lists (“the”) are represented more efficiently using bitvectors of size N
- ▶ State-of-the-art implementations use SIMD Instructions and bit-parallelism to increase decoding speed

Postings List Compression - Fast Searching

- ▶ Often decompressing the whole list is unnecessary
- ▶ Allow of efficient searching in the list is an important operation

Operation $GEQ(x)$ - Find num in list that is Greater or Equal to x

Example:

List: [1,2,5,9,12,15]; $GEQ(6) = 9$; $GEQ(12) = 12$.

Postings List Compression - Fast Searching (GEQ)

- ▶ Compress postings list in blocks of 128 integers at a time
- ▶ For each block store an uncompressed sample value representing the largest (or smallest) value in the block
- ▶ Use sample values to efficiently seek to any position in the postings list without decompressing everything

Operation $GEQ(x)$ algorithm:

- ▶ Binary search over uncompressed sample values to find destination block
- ▶ Decompress destination block to determine final offset in postings list

Postings List Compression - Summary

- ▶ Compress increasing integer sequences
- ▶ Support iterating and searching the compressed sequence
- ▶ Store gaps between adjacent numbers
- ▶ Different compression schemes provide many time-space tradeoffs

Efficient Query Processing (Top- k Retrieval)

IDEA:

Retrieve the top k items for a given query without having to evaluate all documents.

- ▶ Web search engines return only the top-10 results to the user
- ▶ For most queries most users generally do not retrieve more documents
- ▶ No need to score all possible documents to produce a “complete” ranking

Top- k Retrieval - Concepts

- ▶ Avoid scoring documents we know will not appear in the top- k result list
- ▶ For a given similarity metric (for example BM25), prestore some information for each term to avoid scoring
- ▶ Incorporate with block based compression schemes for efficient query processing
- ▶ Utilise the $\text{GEQ}(x)$ operation to avoid decompression of large parts of postings lists

Top- k Retrieval - Review BM25

BM25 computation for one document:

$$S_{Q,d}^{\text{BM25}} = \sum_{t \in Q} \frac{(k_1 + 1)f_{d,t}}{k_1 \left(1 - b + b \frac{L_d}{L_{\text{avg}}}\right) + f_{d,t}} \cdot \ln \left(\frac{N - f_t + 0.5}{f_t + 0.5} \right)$$

Symbol	Description
Q	Multi set of query terms
t	Query term in Q
d	Evaluated document
$f_{d,t}$	Frequency of t in d (Term frequency)
N	Number of documents in the collection
f_t	Number of documents containing t (Document frequency)
L_d	Length of document d
L_{avg}	Average document length in the collection
k_1, b	Constants

Top- k Retrieval - Inefficient Evaluation of BM25

BM25 computation for one document:

$$S_{Q,d}^{\text{BM25}} = \sum_{t \in Q} \underbrace{\frac{(k_1 + 1)f_{d,t}}{k_1 \left(1 - b + b \frac{L_d}{L_{\text{avg}}}\right) + f_{d,t}}}_{=TF} \cdot \underbrace{\ln \left(\frac{N - f_t + 0.5}{f_t + 0.5} \right)}_{=IDF}$$

Inefficient Evaluation:

- ▶ For each t in Q compute IDF in $O(|Q|)$ time
- ▶ For each d in the document collection containing any t in Q evaluate TF (Potentially $O(N)$ time!)
- ▶ Return the top- k highest scoring documents

Top- k Retrieval - Evaluation with WAND

Basic Idea:

- ▶ Keep track of the top- k highest scoring documents
- ▶ For each unique term in the collection store the **maximum contribution** it can have to *any* document score in the collection
- ▶ Skip over documents that can not enter the top- k ranked result list

Citation: Andrei Z. Broder, David Carmel, Michael Herscovici, Aya Soffer, Jason Y. Zien: Efficient query evaluation using a two-level retrieval process. CIKM: 426-434 (2003)

WAND - Maximum contribution of a term t

Maximum contribution

The **Maximum contribution** of a term t is the largest score **any** document in the collection can have for the query Q_t only consisting of t .

- ▶ Depends on the similarity measure
- ▶ Can be computed at construction time of the index
- ▶ Only requires storing a single floating point number for each list
- ▶ Can be used to **overestimate** the score of a document in a multi term query

WAND - Example

Query Q : The quick brown fox

with $k = 2$

The

2	3	7	8	9	10	11	12	13	17	18	19	...
---	---	---	---	---	----	----	----	----	----	----	----	-----

quick

5	6	9	11	14	18
---	---	---	----	----	----

brown

2	4	5	15	42	84	96
---	---	---	----	----	----	----

fox


5	7	8	13
---	---	---	----

WAND - Example

Query Q : The quick brown fox

with $k = 2$


Maximum Contribution for each query term



WAND - Example

Query Q : The quick brown fox

with $k = 2$




WAND - Example

Query Q : The quick brown fox

with $k = 2$

(1) Reorder based on current id and score smallest.




WAND - Example

Query Q : The quick brown fox

with $k = 2$

(2) Move pointer of scored elements.




WAND - Example

Query Q : The quick brown fox

with $k = 2$

(3) Reorder based on current id and score smallest.




WAND - Example

Query Q : The quick brown fox

with $k = 2$

(4) Move pointer of scored elements




WAND - Example

Query Q : The quick brown fox

with $k = 2$

(5) Reorder based on current id and score smallest.




WAND - Example

Query Q : The quick brown fox

with $k = 2$

(6) Decide if we need to score smallest id.




WAND - Example

Query Q : The quick brown fox

with $k = 2$

(7) Replace 3 with 4 on the heap.




WAND - Example

Query Q : The quick brown fox

with $k = 2$

(8) Sort by current id. Evaluate 5. Add to Heap.




WAND - Example

Query Q : The quick brown fox

with $k = 2$

(9) Move pointers and sort.




WAND - Example

Query Q : The quick brown fox

with $k = 2$

(10) Use max to skip scoring smallest.




WAND - Example - Fast Forward

Query Q : The quick brown fox

with $k = 2$

Do we have to evaluate document 9?




WAND - Example - Fast Forward

Query Q : The quick brown fox

with $k = 2$

Do we have to evaluate document 9?




WAND - Example - Fast Forward

Query Q : The quick brown fox

with $k = 2$

What is the next document that has to be evaluated?




WAND - Example - Fast Forward

Query Q : The quick brown fox

with $k = 2$

What is the next document that has to be evaluated?

Max




13, as $0.9 + 1.9 + 7.1 > 6.3$!

WAND - Example - Fast Forward

Query Q : The quick brown fox

with $k = 2$


Fast forward smaller ids to 13 (GEQ) and sort.



WAND - Example - Evaluate 13?

Query Q : The quick brown fox

with $k = 2$



WAND - Algorithm

Given Q , k and the postings lists $L[0 \dots |Q| - 1]$ with:

$L[i].max \triangleq$ the **maximum contribution** of the list

$L[i].cur \triangleq$ the current element of the list

```
1: function WAND( $Q, k, L[0 \dots |Q| - 1]$ )
2:    $TopDocs = \emptyset$  ▷ Min Heap of size  $k$ 
3:   Threshold  $\Theta = 0$  ▷ Smallest score in  $TopDocs$ 
4:   while Not all lists are processed do
5:     Sort  $L$  based on  $L[i].cur$ 
6:     Select pivot list  $p$  such that  $\sum_0^{p-1} L[i].max \geq \Theta$ 
7:     Forward all lists  $L[0 \dots |p| - 1]$  to  $d_p = L[p].cur$ 
8:     Compute  $S_{Q, d_p}^{BM25}$  and insert into  $TopDocs$  if score  $> \Theta$ 
9:      $\Theta = \min(TopDocs)$  or 0 if  $|TopDocs| < k$ .
10:  end while
11:  Return  $TopDocs$ 
12: end function
```

WAND - Discussion

- ▶ Use max contribution of query term to overestimate score of a document.
- ▶ Do not score document if it can not enter the top- k heap.
- ▶ Utilize GEQ function of compressed representation to skip over large parts of the postings lists.
- ▶ Similarity metric fixed at index construction time.
- ▶ Works very well in practice.

WAND - Performance

Method	$k = 10$		$k = 1,000$	
	Avg.	Med.	Avg.	Med.
Exhaustive	100.0	100.0	100.0	100.0
WAND	3.5	1.0	28.0	11.7

Figure: Fraction of pointers processed as a percentage of the total number of pointers associated with each query, GOV2, using TREC topics 701–850. Across the set of queries, the average number of postings per query for exhaustive processing is 1,460,562, and the median number of postings is 1,080,008. The percentages shown in the table are relative to these two numbers.

WAND - Performance II




Figure: Evaluation of one query “north korean counterfeiting” for $k = 10$.

Further Reading

Reading:

- ▶ Manning, Christopher D; Raghavan, Prabhakar; Schütze, Hinrich; Introduction to information retrieval, Cambridge University Press 2008. (Chapter 5)

Additional References:

- ▶ Daniel Lemire, Leonid Boytsov: Decoding billions of integers per second through vectorization. Softw., Pract. Exper. 45(1): 1-29 (2015)
- ▶ Andrei Z. Broder, David Carmel, Michael Herscovici, Aya Soffer, Jason Y. Zien: Efficient query evaluation using a two-level retrieval process. CIKM: 426-434 (2003)